

Python/Numpy for High-Performance Numerical Processing

Jim Pivarski


Princeton University

November 15, 2018

Why Python?





PYPL Popularity of Programming Language





<http://pypl.github.io/PYPL.html>


Why Python in science?



Why Python in science?




Why Python in science?



All of the machine learning libraries I could find either have a Python interface or are primarily/exclusively Python.



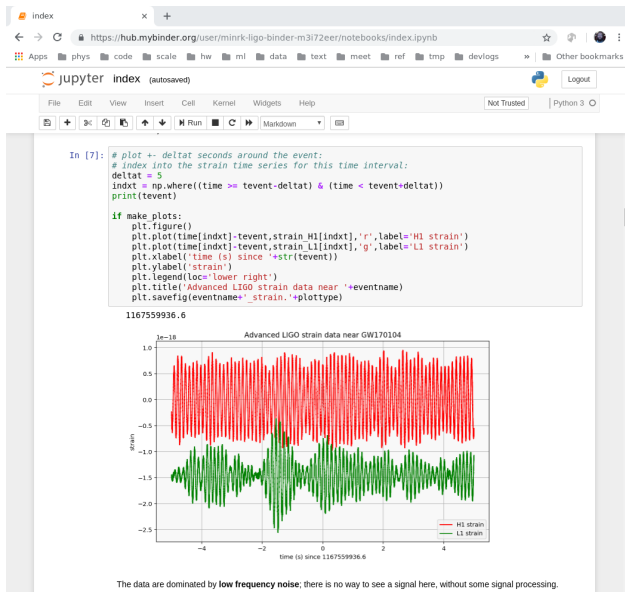
Mentions of Software in Astronomy Publications:



Compiled from NASA ADS [\(code\)](#).


Thanks to Juan Nunez-Iglesias,
Thomas P. Robitaille, and Chris Beaumont.

Why Python in science?







Python's Scientific Stack




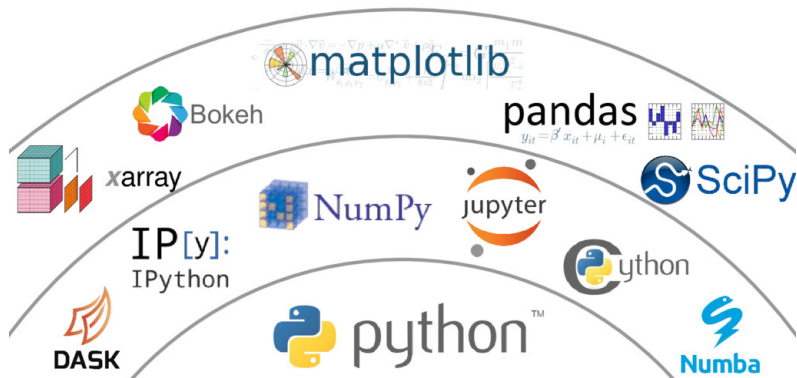


Python's Scientific Stack







Python's Scientific Stack







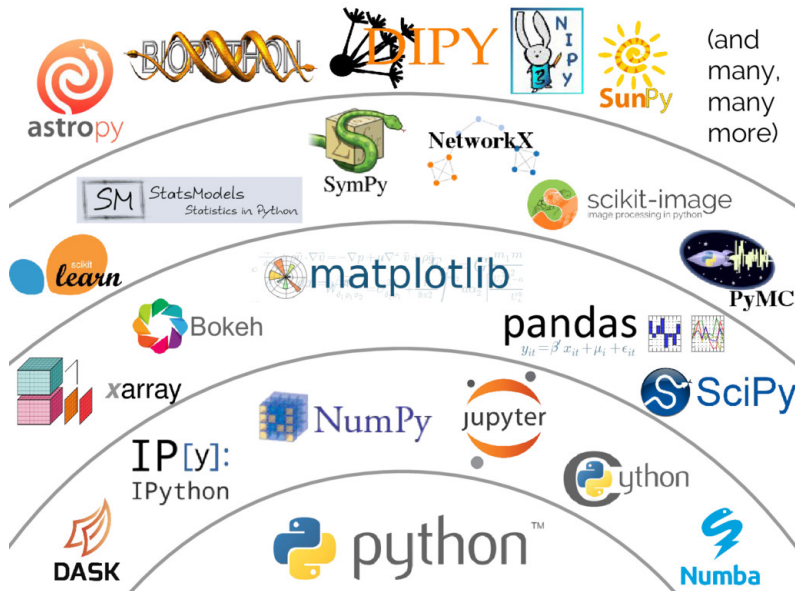
Python's Scientific Stack




Stealing from Jake VanderPlas's *Unexpected Effectiveness* talk




Stealing from Jake VanderPlas's *Unexpected Effectiveness* talk



(and many, many more)



If you're used to writing your own code, searching for tools is eye-opening: you learn what's unique about what you do and what isn't.





In science, we often have to scale up analyses to large datasets.



In science, we often have to scale up analyses to large datasets.

10% faster doesn't mean much, but the difference between
“five minutes” and “overnight” is life-changing.



In science, we often have to scale up analyses to large datasets.

10% faster doesn't mean much, but the difference between
“five minutes” and “overnight” is life-changing.

That's the scale we're talking about between C and Python.



In science, we often have to scale up analyses to large datasets.

10% faster doesn't mean much, but the difference between
“five minutes” and “overnight” is life-changing.

That's the scale we're talking about between C and Python.

But we also need the interactivity of a dynamic language to *develop* the analysis.
(“If we knew what we were doing, it wouldn't be called research.”)



Drive to the airport by car, then take a plane.

Small-scale *project organization* in Python, ignoring performance entirely.

Run over *big data* in compiled code, tuning performance until it no longer matters.

Fit function

[illegible]

PyMinuit

SEAL-MINUIT

```
#include "Unitist/Unitist.h"
#include "Unitist/FunctionMinimax.h"
#include "Unitist/FunctionCross.h"
#include "Unitist/FunctionCross.h"
#include "Unitist/FunctionCross.h"

std::pair<double, double> Minimax::operator()(unsigned int par, unsigned
int maxcalls) const
{
    Minimax mmax = *this(par, maxcalls);
    return mmax();
}

double Minimax::lower(unsigned int par, unsigned int maxcalls) const {
    MinimaxParameterState ugr = Minimax::userState();
    double err = Minimax::userState().err(par);
    Minimax mmax = *this(par, maxcalls);

    double lower = ugr.isval(); if (!err || (1 + ugr.sopt.value()) >
    (mmax.atval() / ugr.parameter(par).lowerlimit()) : ugr.value(par));

    return lower;
}
```

Fitting script

```

get_runs has been given a thorough look-over: it is correct (7 votes)
# get_runs contains all corrections, from numbers of events to
# live time correction.)

from numpy import *
execfile('home/hccann/antitheater/utilities.py')
import glob
import globfn
import sys

def dofiguins(h):
    def gauss(x, n):
        return exp(-(x-0.5)**2/2.0)/sqrt(2.0*np.pi)
    def figuins(n):
        c = 0
        for n in h.data:
            c = log(gauss(n, n))
        return c
    n = Minus(figuins, start=0, l=1, up=0.5)
    n.mplab(0)
    n.mplab(0.1)
    errors = n.minus_errors[0][1] - n.minus_errors[0][0]/0.2
    err = n.minus_errors[0][1] - n.minus_errors[0][0]
    return figuins(0), err, n.mvalue[1], err, lambda x:
    0.2*exp(-(x-n.data[0]-n.value[0])**2/(n.mvalue[1]-n.value[0]))

```

GNU plotutils



Numeric




Plotting script

```
from math import *
import biggles, Numeric, cPickle as pickle
import gbkvif
import gbkvftau

allthat = pickle.load(file('/home/ccarron/antithesis/novemberdata.p'))
u1runs = allthat["u1runs"]
u2runs = allthat["u2runs"]
u3runs = allthat["u3runs"]
...
q = biggles.FramedPlot()
adddata(q, [None], [u1start[u1gbl"], 0.]
addfunc(q, [None], [u2start[u2gbl"], 0])
addfunc(q, [None], [u3start[u3gbl"], 0])
addfunc(q, [None], [None], [0.0000, 0.0000, linewidth='dashed'])
```

(pickle)



Which got me involved in open source (PyMinuit is now “iminuit”)



[What's new?](#) | [Help](#) | [Directory](#) | [Sign in](#)



pyminuit

Minuit numerical function minimization in Python

[Project Home](#) [Downloads](#) [Wiki](#) [Issues](#) [Source](#)

PyMinuit

Minuit numerical function minimization in Python

Minuit

Minuit has been the standard package for minimizing general N-dimensional functions in high-energy physics since its introduction in 1972. It features a robust set of algorithms for optimizing the search, correcting mistakes, and measuring non-linear error bounds. It is the minimization engine used behind-the-scenes in most high-energy physics curve fitting applications.

New: more robust [installation instructions](#)!

Python interface

PyMinuit is an extension module for Python that passes low-level Minuit functionality to Python functions. Interaction and data exploration is more user-friendly, in the sense that the user is protected from segmentation faults and index errors, parameters are referenced by their names, even in correlation matrices, and Python exceptions can be passed from the objective function during the minimization process. This extension module also makes it easier to calculate Minos errors and contour curves at an arbitrary number of sigmas from the minimum, and features a new N-dimensional scanning utility.

Search Projects

Search the Web

License: [GNU General Public License v2](#)

Labels: [python](#), [minuit](#), [optimization](#), [computation](#), [HEP](#), [math](#), [fitting](#), [physics](#)

Featured Downloads:

[Show all](#)

↓ [Minuit-1.7.9.tar.gz](#)

↓ [pyminuit-1.0.2.tgz](#)

Featured Wiki Pages:

[Show all](#)

[FunctionReference](#)

[GettingStartedGuide](#)

[HowToInstall](#)

Links:

- [Official Minuit homepage at CERN](#)
- [Minuit documentation through the ages](#)

Project owners:

[Join project](#)

[jpivarski](#)



The key to ecosystem development was a common array library

- 1994 **Python** 1.0 released.
- 1995 First array package: **Numeric** (a.k.a. Numerical, Numerical Python, NumPy).
- 2001 Diverse scientific codebases merged into **SciPy**.
- 2003 **Matplotlib**
- 2003 Numeric was limited; **numarray** appeared as a competitor with more features (memory-mapped files, alignment, record arrays).
- 2005 Two packages were incompatible; could not integrate numarray-based code into SciPy. Travis Oliphant merged the codebases as **Numpy**.
- 2008 **Pandas**
- 2010 **Scikit-Learn**
- 2011 **AstroPy**
- 2012 **Anaconda**
- 2014 **Jupyter**
- 2015 **Keras**



The key to ecosystem development was a common array library

- 1994 **Python** 1.0 released.
- 1995 First array package: **Numeric** (a.k.a. Numerical, Numerical Python, NumPy).
- 2001 Diverse scientific codebases merged into **SciPy**.
- 2003 **Matplotlib**
- 2003 Numeric was limited; **numarray** appeared as a competitor with more features (memory-mapped files, alignment, record arrays).
- 2005 Two packages were incompatible; could not integrate numarray-based code into SciPy. Travis Oliphant merged the codebases as **Numpy**.
- 2008 **Pandas**
- 2010 **Scikit-Learn**
- 2011 **AstroPy**
- 2012 **Anaconda**
- 2014 **Jupyter**
- 2015 **Keras**

The scientific Python ecosystem could have failed before it started if the Numeric/numarray split hadn't been resolved!

Numpy is high-level, array-at-a-time math



```
>>> import numpy
>>> a = numpy.arange(12)
>>> a
array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11])
>>> a.shape = (3, 4)
>>> a
array([[ 0,  1,  2,  3],
       [ 4,  5,  6,  7],
       [ 8,  9, 10, 11]])
>>> a.sum(axis=0)
array([12, 15, 18, 21])
>>> a.min(axis=1)
array([0, 4, 8])
>>> a**2
array([[ 0,  1,  4,  9],
       [16, 25, 36, 49],
       [64, 81, 100, 121]])
>>> numpy.sqrt(a)
array([[0.          ,  1.          ,  1.41421356,  1.73205081],
       [2.          ,  2.23606798,  2.44948974,  2.64575131],
       [2.82842712,  3.          ,  3.16227766,  3.31662479]])
```



Although you can write Python `for` loops over Numpy arrays, you don't reap the benefit unless you express your calculation in Numpy universal functions (ufuncs).

```
pz = numpy.empty(len(pt))  
for i in range(len(pt)):  
    pz[i] = pt[i]*numpy.sinh(eta[i])
```

VS `pz = pt * numpy.sinh(eta)`

$\mathcal{O}(N)$ Python bytecode instructions, type-checks, interpreter locks.

$\mathcal{O}(1)$ Python bytecode instructions, type-checks, interpreter locks.

$\mathcal{O}(N)$ statically typed, probably vectorized native bytecode operations on contiguous memory.



Although you can write Python `for` loops over Numpy arrays, you don't reap the benefit unless you express your calculation in Numpy universal functions (ufuncs).

```
pz = numpy.empty(len(pt))  
for i in range(len(pt)):  
    pz[i] = pt[i]*numpy.sinh(eta[i])
```

VS `pz = pt * numpy.sinh(eta)`

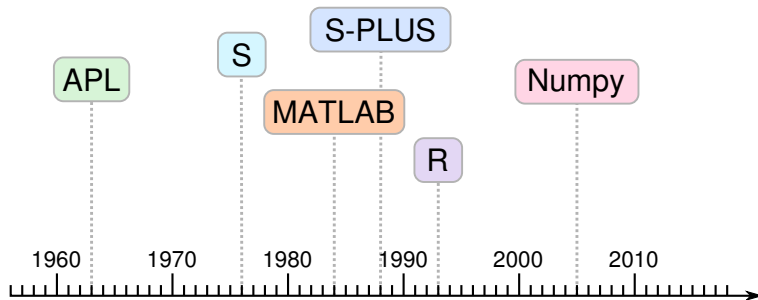
$\mathcal{O}(N)$ Python bytecode instructions, type-checks, interpreter locks.

$\mathcal{O}(1)$ Python bytecode instructions, type-checks, interpreter locks.

$\mathcal{O}(N)$ statically typed, probably vectorized native bytecode operations on contiguous memory.


In other words, a Single (Python) Instruction on Multiple Data.
Conceptually similar to SIMD, the program flow of GPUs.

APL, “A Programming Language” introduced the idea of single commands having sweeping effects across large arrays.



All members of the APL family are intended for interactive data analysis. Numpy, however, is a library in a general-purpose language, not a language in itself.

APL pioneered conciseness;
discovered the mistake of being too concise.



Conway's Game of Life was one line of code:

$$\text{life} \leftarrow \{\uparrow 1 \quad \omega \vee . \wedge 3 \quad 4 = +/,^{\neg} 1 \quad 0 \quad 1 \circ . \Theta^{\neg} 1 \quad 0 \quad 1 \circ . \Phi \subset \omega\}$$


“Map” was implicit, “reduce” was a slash, functions were symbols. For example:

APL	Numpy
$m \leftarrow +/(3 + \iota 4)$	<code>m = (numpy.arange(4) + 3).sum()</code>



As an array abstraction, Numpy presents a high-level way for users to think about vectorization.

Vectorization is key to using GPUs and modern CPUs efficiently.



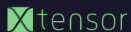
HIGH PERFORMANCE WITH CUDA

CuPy is an open-source matrix library accelerated with NVIDIA CUDA. It also uses CUDA-related libraries including cuBLAS, cuDNN, cuRand, cuSolver, cuSPARSE, cuFFT and NCCL to make full use of the GPU architecture.

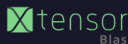
PROJECTS

QuantStack developers contribute to a number of open-source projects, including jupyter, xtensor, bqplot, conda-forge and many others.

High-Performance-Computing



C++ tensor algebra library



BLAS extension to xtensor



C++ wrappers for SIMD intrinsics
and optimized math
implementations



Plan for the day



10:00	Intro talk	1-intro.pdf
11:00	Just Numpy	2-just-numpy.ipynb
12:00	Lunch	
1:00	Numpy ecosystem talk	3-ecosystem.pdf
2:00	Pandas	4-pandas.ipynb
	Dask & multiprocessing	5-dask.ipynb
	Coffee break	
3:00	Numba, Cython, pybind11	6-compilers.ipynb
	CuPy, Numba-GPU, PyCUDA	7-gpu.ipynb
4:00	ctypes & low-level hackery	8-low-level.ipynb

Skills-based Numpy tutorial with a couple of exercises in the morning: how to think in SIMD.

Overview of libraries in the afternoon: where to look for solutions to your problems.